

OpenCL exercise 4: Matrix Multiplication

Kaicong Sun

Local Memory

- ▶ Local memory: Shared by all work items of one work group
 - ▶ Two work items in the same work group will see the same data
 - ▶ Two work items in the different work groups will see different data
- ▶ Size: 16kB or 48kB
- ▶ Significantly higher memory bandwidth compared to global memory
- ▶ Significantly lower latency compared to global memory
 - ▶ Random accesses patterns are fast
- ▶ Used for:
 - ▶ Manually caching data from global memory
 - ▶ Storage for data being worked on
- ▶ In Cuda: “Shared memory”
 - ▶ Has nothing to do with Cuda “Local memory”

Local Memory Allocation

After compiling the program, the local memory size can not be modified. In other words, the size must be known before compiling. On the kernel side:

- ▶ Declare “i” as 32-bit integer in local memory

```
__local int i;
```

- ▶ Declare “a” as 2D array of unsigned 32-bit integers with a size of 10x15

```
__local uint a[10][15];
```

- ▶ “10” and “15” must be compile-time constants
- ▶ Wait until all threads have reached this point and prevent any local memory access from being moved across this line

```
barrier(CLK_LOCAL_MEM_FENCE);
```

Dynamic Local Memory Allocation

The size of the local memory is not necessary to be known before compiling. On the kernel side:

- ▶ Kernel definition:

```
__kernel void kernel1(__local float* localMem) { ... }
```

- ▶ Use inside kernel

```
localMem[i] = 10;  
foo = localMem[j];
```

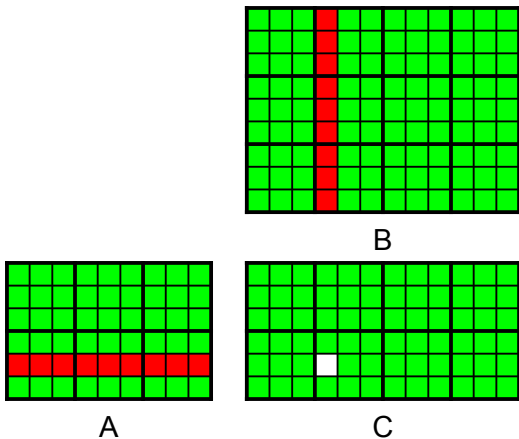
- ▶ Calling the kernel

```
kernel1.setArg(0, cl::Local(2*wgX*wgY*sizeof(float)));  
queue.enqueueNDRangeKernel(kernel1, ...);
```

- ▶ `localMem[]` will contain space for `wgX * wgY` floats

Matrix Multiplication

- ▶ 3 matrices: **A**, **B**, **C**, calculate $C = AB$
- ▶ $C_{i,j} = \sum_{k=1}^n A_{i,k} B_{k,j}$

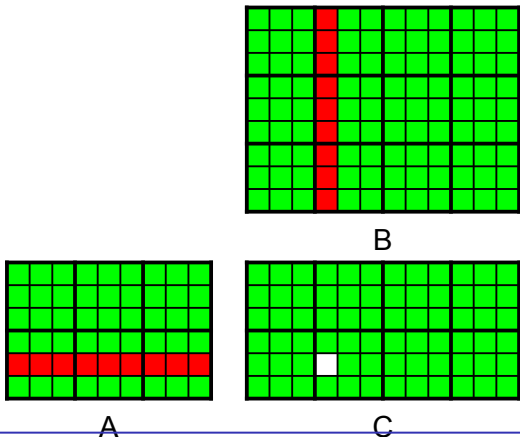


Host code

```
1  for (std::size_t j = 0; j < countAY; j++) {  
2      for (std::size_t i = 0; i < countBX; i++) {  
3          float sum = 0;  
4          for (std::size_t k = 0; k < countAX_BY; k++) {  
5              float a = h_inputA[k + j * countAX_BY];  
6              float b = h_inputB[i + k * countBX];  
7              sum += a * b;  
8          }  
9          h_outputC[i + j * countBX] = sum;  
10     }  
11 }
```

Task 1: Simple GPU implementation

Implement host code on GPU, use one work item for each element in C (the result matrix).




Task 2: Local memory


Take advantage of local memory to speed up the calculation

- ▶ Input Matrices are split into quadratic blocks, size of block = size of work group
- ▶ Each work group will:
 - ▶ For each input block needed:
 - ▶ Load the block from global input buffer *A* into local memory
 - ▶ Load the block from global input buffer *B* into local memory
 - ▶ Barrier
 - ▶ Use the data from local memory to calculate the result
 - ▶ Barrier
 - ▶ Store the result


Step 1



B




A




C


Step 2



B




A

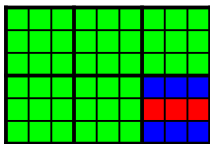


C

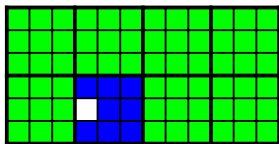
Step 3



B



A



C

Pseudocode

```
1  for (uint j = 0; j < countAY; j++) {
2      for (uint i = 0; i < countBX; i++) {
3          float sum = 0;
4          int k = get_local_id(0);
5          int g = get_local_id(1);
6          __local float l_A[WG_SIZE][WG_SIZE];
7          __local float l_B[WG_SIZE][WG_SIZE];
8          // loop over the submatrices
9          for (uint bs = 0; bs < countAX_BY; bs += WG_SIZE) {
10             //Copy blocks of d_inputA, d_inputB to local memory
11             l_A[g][k] = d_inputA[(k+bs) + j * countAX_BY];
12             l_B[g][k] = d_inputB[i + (g+bs) * countBX];
13
14             barrier(CLK_LOCAL_MEM_FENCE);
15             for (uint m = 0; m < WG_SIZE; m++)
16                 sum += l_A[g][m] * l_B[m][k];
17             barrier(CLK_LOCAL_MEM_FENCE);
18         }
19         d_outputC[i + j * countBX] = sum;
20     }
21 }
```

Hints

- ▶ For copying a block of data (with size of block = size of work group) into local memory, each work item has to copy one value
 - ▶ Two consecutive work items in X-direction should read two consecutive values from global memory
 - ▶ `localMem[get_local_id(1)][get_local_id(0)] = globalMem[x + width * y];`
- ▶ To make sure a kernel is only called with a certain work group size:

```
__attribute__((reqd_work_group_size(10, 20, 1)))  
__kernel void kernel1() { ... }
```

Kernel can only be called with a 2D work group with 10x20 elements.

Task 3/4

Two (optional) tasks:

- ▶ Task 3: Make a copy of the kernel of Task 2 and modify it so that it doesn't need any compile-time knowledge of the work group size (i.e. `WG_SIZE`)
 - ▶ Will need a dynamically sized local memory area
 - ▶ Will be slower than kernel 2

```
matrixMulKernel.setArg(4, cl::Local(2 * wgSize
* wgSize * sizeof(float)));
__local float* l_A = localMem;
__local float* l_B = localMem + get_local_size(0)
* get_local_size(1);
```

- ▶ Task 4: Create a kernel similar to kernel 1, but use OpenCL images for *A* and *B*
 - ▶ Will be slower than kernel 2 and kernel 3
 - ▶ Might be slower or faster than kernel 1